

**VIT<sup>®</sup>**

**Vellore Institute of Technology**  
(Deemed to be University under section 3 of UGC Act, 1956)

A Project Report

submitted as part of the course

**Information Visualization (CSE3044)**

School of Computer Science and Engineering

VIT Chennai

Winter Semester 2020-2021

**Course Faculty: Dr. Arun Kumar Sivaraman**

***PROJECT:***

**CAMPUS RECRUITMENT**

*NAME: SANJANA ALAHAM*

*REG NO: 19MIA1055*

## **ABSTRACT**

Campus recruitment is a strategy for sourcing, engaging and hiring young talent for internship and entry-level positions. College recruiting is typically a tactic for medium- to large-sized companies with high-volume recruiting needs, but can range from small efforts (like working with university career centers to source potential candidates) to large-scale operations (like visiting a wide array of colleges and attending recruiting events throughout the spring and fall semester). Campus recruitment often involves working with university career services centers and attending career fairs to meet in-person with college students and recent graduates.

## **INTRODUCTION**

This data set consists of Placement data of students . It includes secondary and higher secondary school percentage and specialization. It also includes degree specialization, type and Work experience and salary offers to the placed students.

### **Questions:**

1. Which factor influenced a candidate in getting placed?
2. Does percentage matter for one to get placed?
3. Which degree specialization is much demanded by corporations?
4. Does work experience matter for one to get placed?
5. Play with the data conducting all statistical tests.

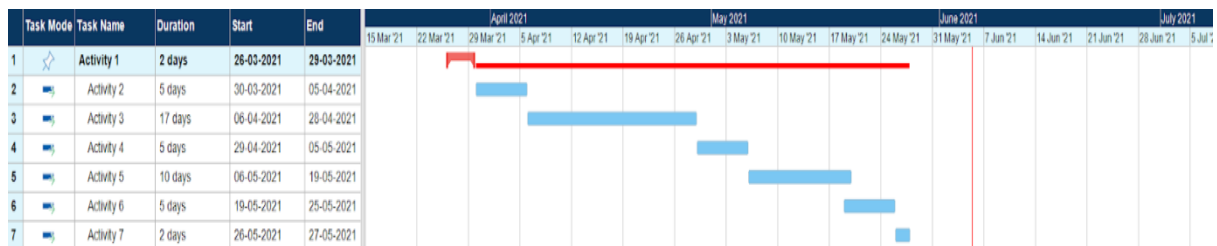
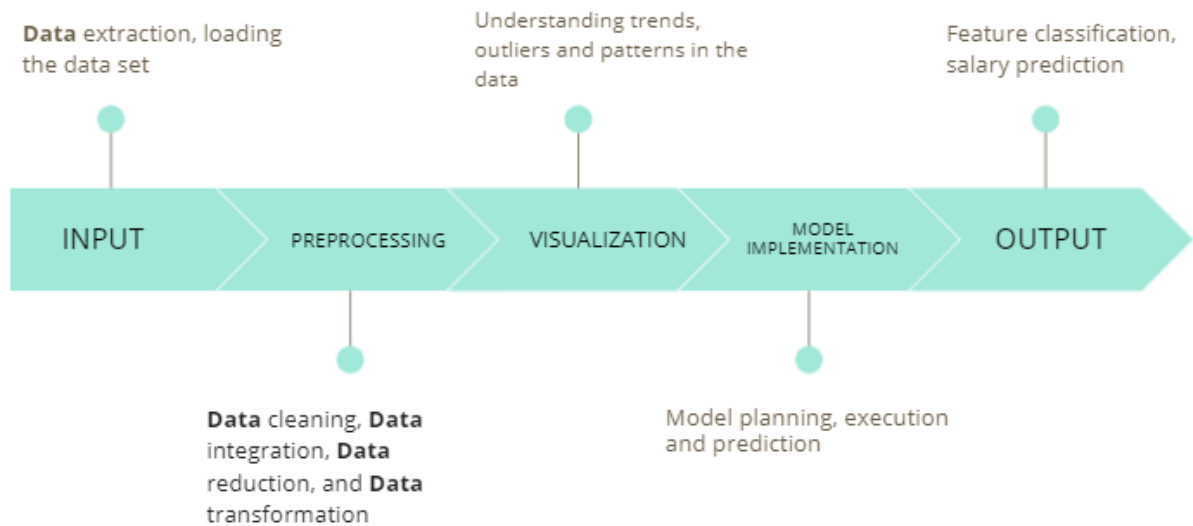
In this project we are going to visualize the chances and possibility for a candidate getting placed in a company based on the factors given below:

1. Work experience
2. 10th grade percentage
3. 12th Grade percentage
4. Degree percentage
5. Gender
6. Field of Degree

# PROJECT MODULES

## MODULE-1

It is the work flow i.e., the process of our project



## MODULE-2

### Overview of the data set

Overview	Reproduction	Warnings 1
Dataset statistics		
Number of variables	14	
Number of observations	215	
Missing cells	67	
Missing cells (%)	2.2%	
Duplicate rows	0	
Duplicate rows (%)	0.0%	
Total size in memory	117.3 KiB	
Average record size in memory	558.6 B	
Variable types		
CAT	7	
NUM	6	
BOOL	1	

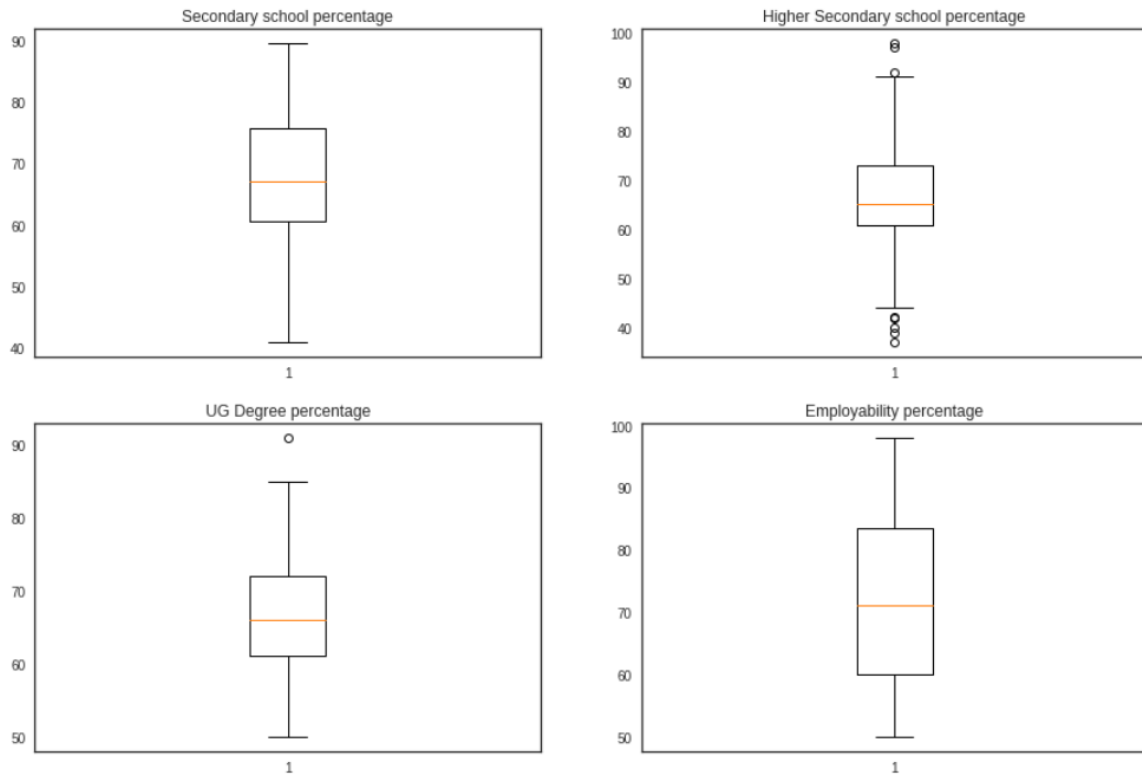
## MODULE-3

- 67 Missing values in Salary for students who didn't get placed. **NaN Value needs to be filled.**
- **Data is not scaled.** Salary column ranges from 200k-940k, rest of numerical columns are percentages.
- 300k at 75th Percentile goes all the way up to 940k max, in Salary (high skewnwss). Thus, **outliers at high salary end.**

### Outlier Detection

```
In [91]: plt.figure(figsize = (15, 10))
plt.style.use('seaborn-white')
ax=plt.subplot(221)
plt.boxplot(data['ssc_p'])
ax.set_title('Secondary school percentage')
ax=plt.subplot(222)
plt.boxplot(data['hsc_p'])
ax.set_title('Higher Secondary school percentage')
ax=plt.subplot(223)
plt.boxplot(data['degree_p'])
ax.set_title('UG Degree percentage')
ax=plt.subplot(224)
plt.boxplot(data['etest_p'])
ax.set_title('Employability percentage')
```

```
Out[91]: Text(0.5, 1.0, 'Employability percentage')
```

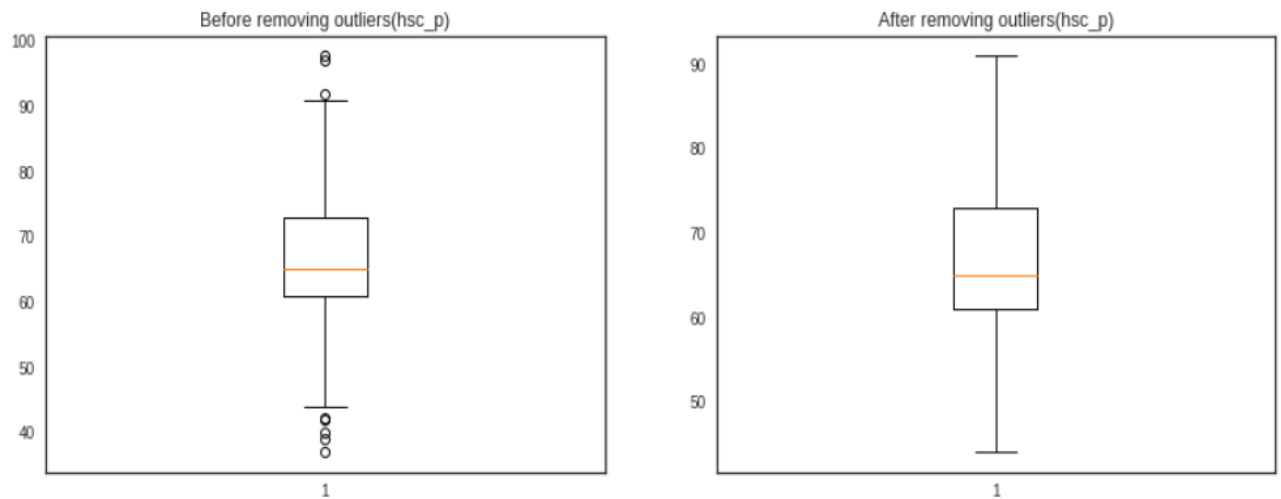


```
In [92]: Q1 = data['hsc_p'].quantile(0.25)
Q3 = data['hsc_p'].quantile(0.75)
IQR = Q3 - Q1    #IQR is interquartile range.

filter = (data['hsc_p'] >= Q1 - 1.5 * IQR) & (data['hsc_p'] <= Q3 + 1.5 * IQR)
data_new=data.loc[filter]
```

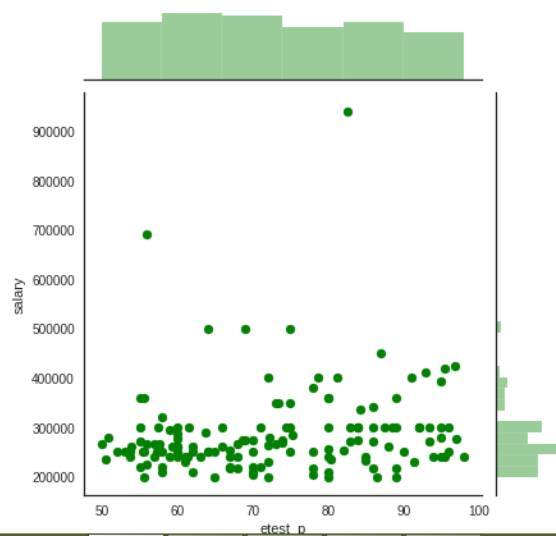
```
In [93]: plt.figure(figsize = (15, 5))
plt.style.use('seaborn-white')
ax=plt.subplot(121)
plt.boxplot(data['hsc_p'])
ax.set_title('Before removing outliers(hsc_p)')
ax=plt.subplot(122)
plt.boxplot(data_new['hsc_p'])
ax.set_title('After removing outliers(hsc_p)')
```

Out[93]: Text(0.5, 1.0, 'After removing outliers(hsc\_p)')



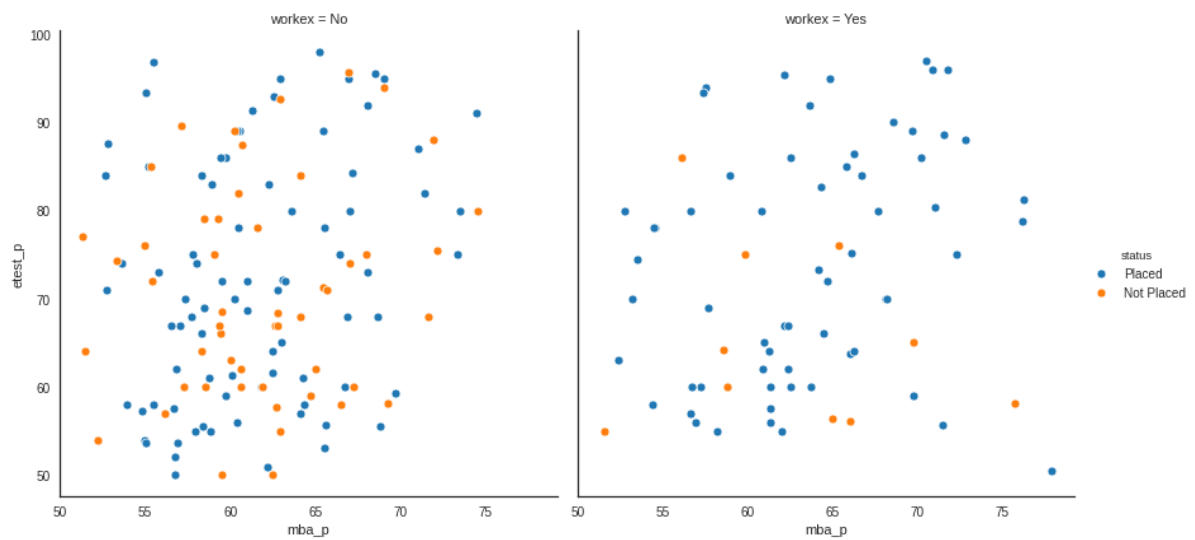
```
In [94]: sns.jointplot(x='etest_p',y='salary',data=data_new,color='green')
```

```
Out[94]: <seaborn.axisgrid.JointGrid at 0x7fbd77a5710>
```



```
In [95]: g=sns.FacetGrid(data=data_new,col='workex',hue='status',height=6)
g.map(sns.scatterplot,'mba_p','etest_p',label='status')
g.add_legend()
```

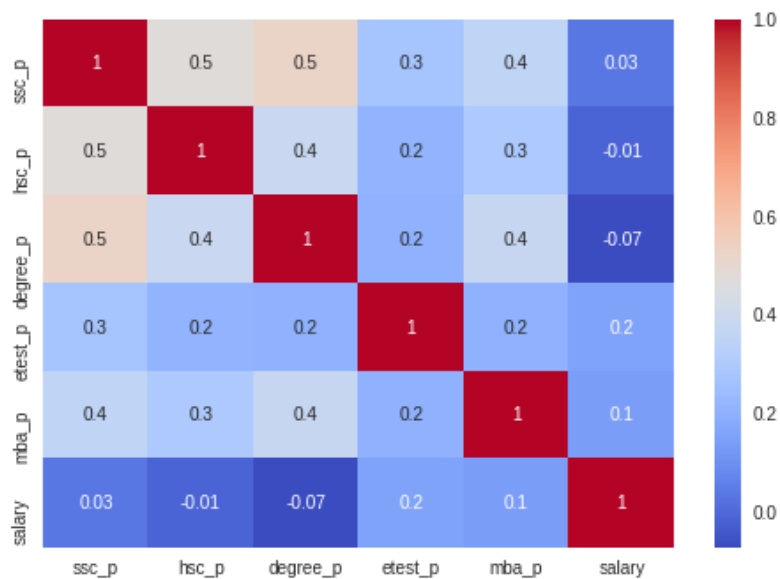
```
Out[95]: <seaborn.axisgrid.FacetGrid at 0x7fbd74bee80>
```



## Coorelation between academic percentages

```
In [96]: p=data_new.corr()
sns.heatmap(p,cmap='coolwarm',annot=True,fmt='.1g')
```

```
Out[96]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbdb7410eb8>
```

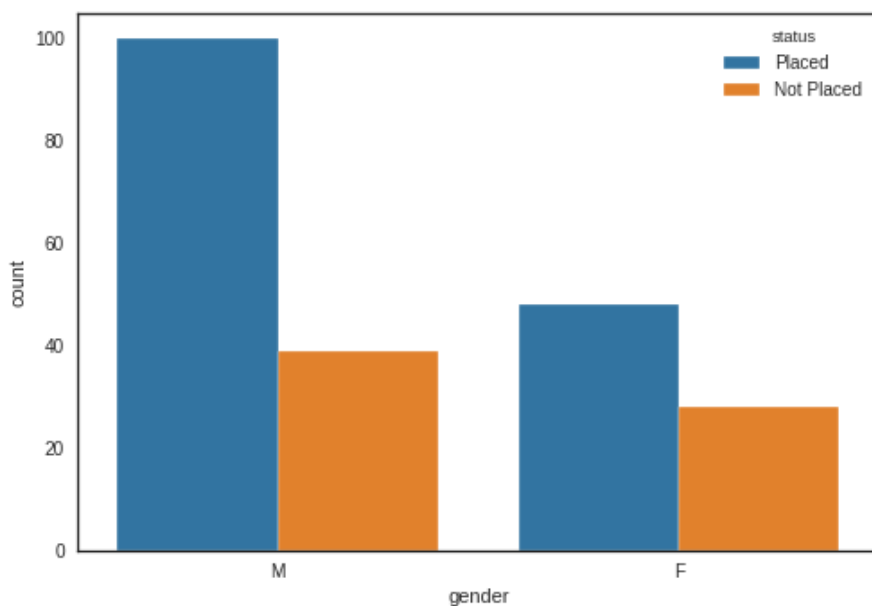


## Vizualizing individual features

### Feature: Gender

Does gender affect placements?

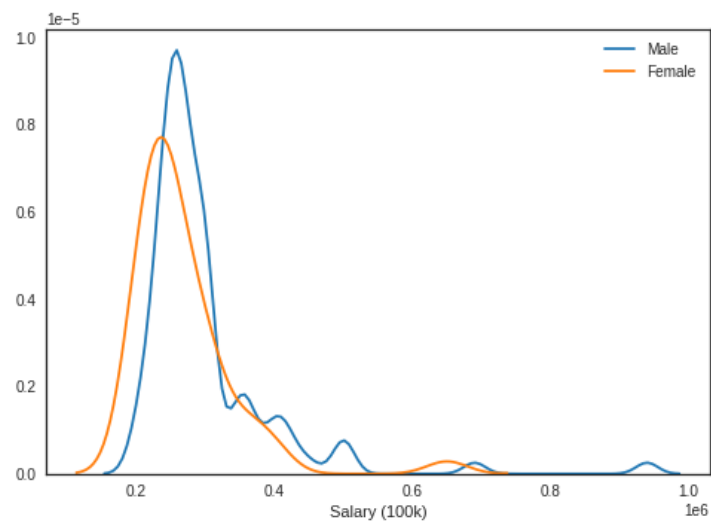
```
1]: data.gender.value_counts()  
   # Almost double  
  
2]: M    139  
   F     76  
   Name: gender, dtype: int64  
  
3]: sns.countplot("gender", hue="status", data=data)  
   plt.show()
```



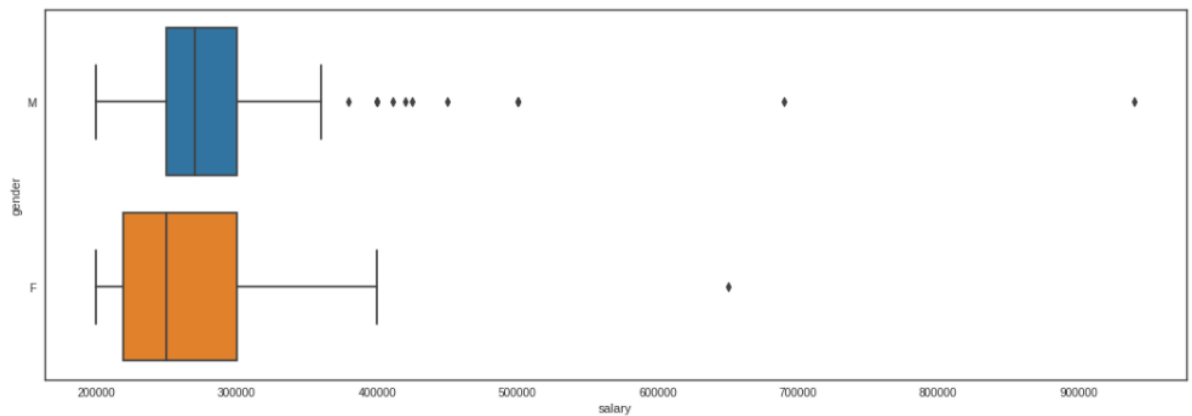
```
4]: #This plot ignores NaN values for salary, igoring students who are not placed  
   sns.kdeplot(data.salary[ data.gender=="M"])  
   sns.kdeplot(data.salary[ data.gender=="F"])  
   plt.legend(["Male", "Female"])  
   plt.xlabel("Salary (100k)")  
   plt.show()
```



```
In [99]: #This plot ignores NaN values for salary, igoring students who are not placed
sns.kdeplot(data.salary[ data.gender=="M"])
sns.kdeplot(data.salary[ data.gender=="F"])
plt.legend(["Male", "Female"])
plt.xlabel("Salary (100k)")
plt.show()
```



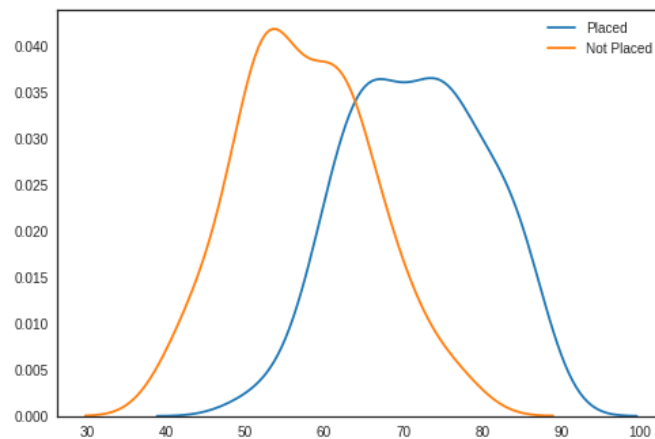
```
In [100]: plt.figure(figsize =(18,6))
sns.boxplot("salary", "gender", data=data)
plt.show()
```



## Feature: SSC\_P (Secondary Education percentage), SSC\_B (Board Of Education)

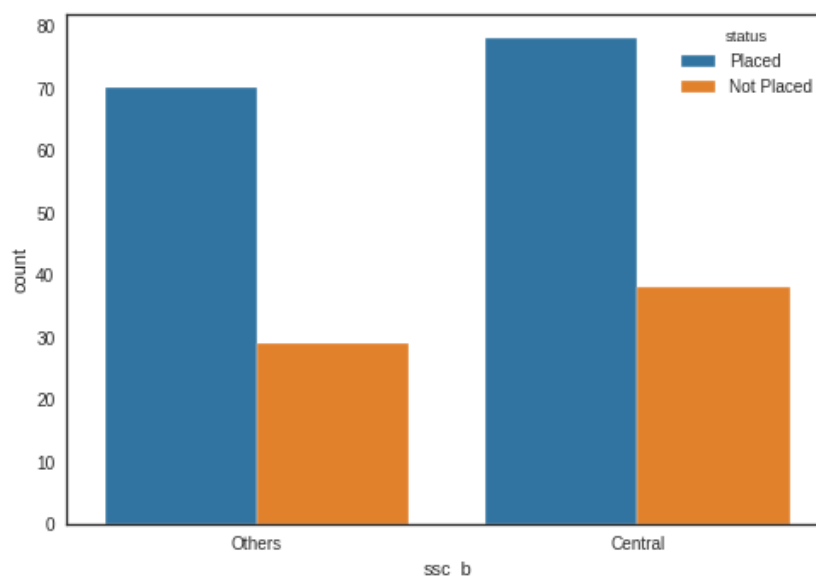
Does Secondary Education affect placements?

```
In [101]: #Kernel-Density Plot
sns.kdeplot(data.ssc_p[ data.status=="Placed"])
sns.kdeplot(data.ssc_p[ data.status=="Not Placed"])
plt.legend(["Placed", "Not Placed"])
plt.xlabel("Secondary Education Percentage")
plt.show()
```



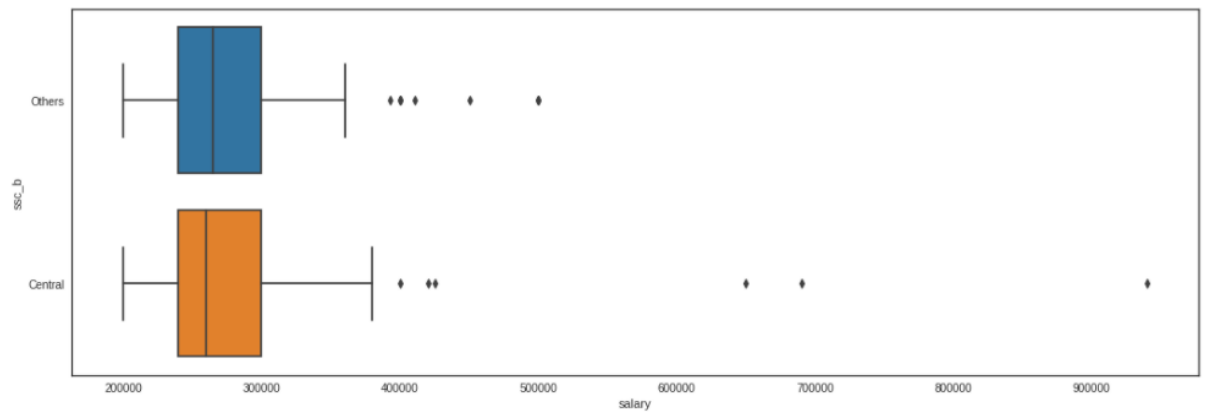
- All students with Secondary Education Percentage above 90% are placed
- All students with Secondary Education Percentage below 50% are not-placed
- **Students with good Secondary Education Percentage are placed on average.**

```
In [102]: sns.countplot("ssc_b", hue="status", data=data)
plt.show()
```



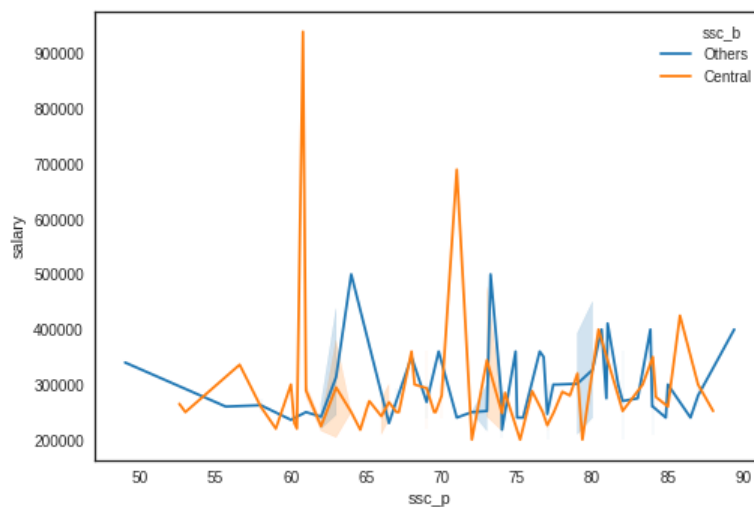
- Board Of Education does not affect Placement Status much

```
In [103]: plt.figure(figsize=(18,6))
sns.boxplot("salary", "ssc_b", data=data)
plt.show()
```



- Outliers on both, but students from Central Board are getting the highly paid jobs.

```
In [104]: sns.lineplot("ssc_p", "salary", hue="ssc_b", data=data)
plt.show()
```

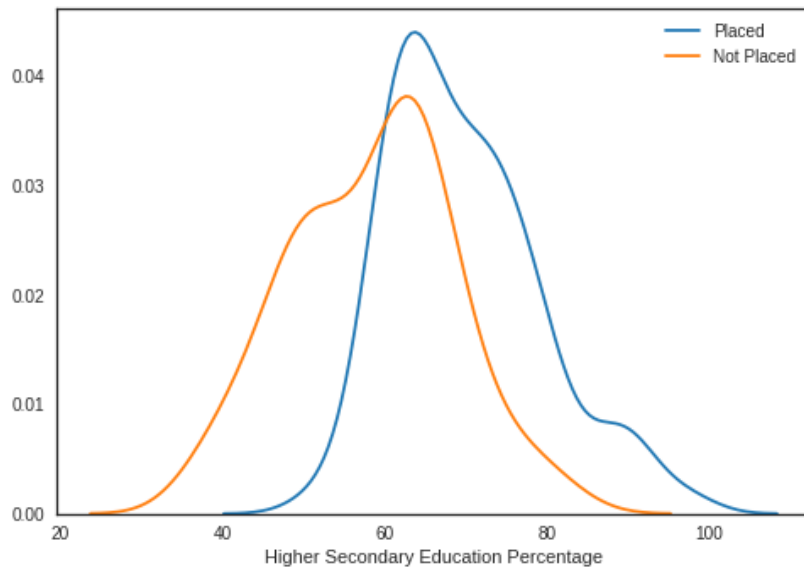


- No specific pattern (correlation) between Secondary Education Percentage and Salary.
- Board of Education is Not Affecting Salary

**Feature: HSC\_P (Higher Secondary Education percentage), HSC\_B (Board Of Education), HSC\_S (Specialization in Higher Secondary Education)**

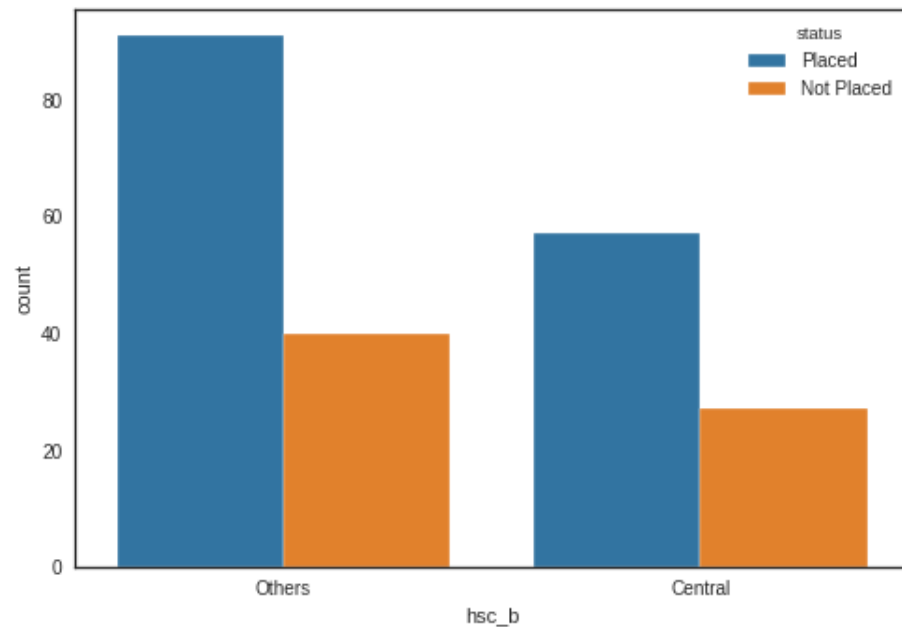
**Does Higher Secondary School affect Placements?**

```
In [105]: #Kernel-Density Plot
sns.kdeplot(data.hsc_p[ data.status=="Placed"])
sns.kdeplot(data.hsc_p[ data.status=="Not Placed"])
plt.legend(["Placed", "Not Placed"])
plt.xlabel("Higher Secondary Education Percentage")
plt.show()
```



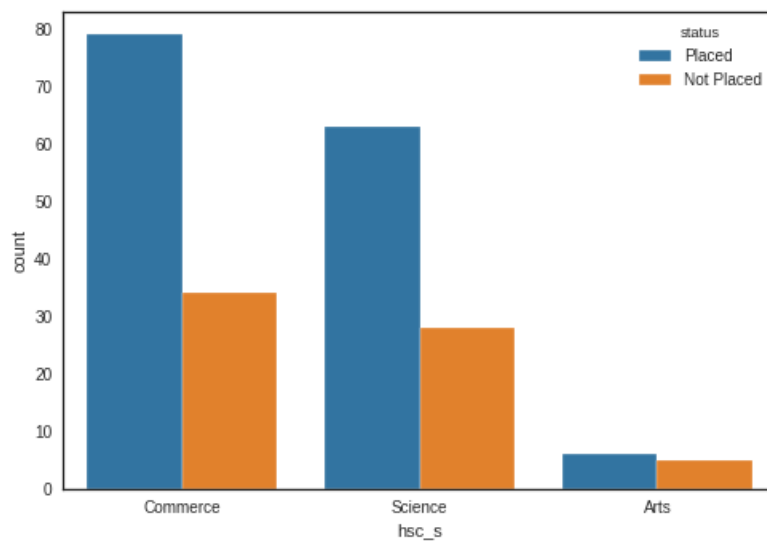
- Overlap here too. More placements for percentage above 65%
- Straight drop below 60 in placements -> Percentage must be atleast 60 for chance of being placed

```
In [106]: sns.countplot("hsc_b", hue="status", data=data)
plt.show()
```



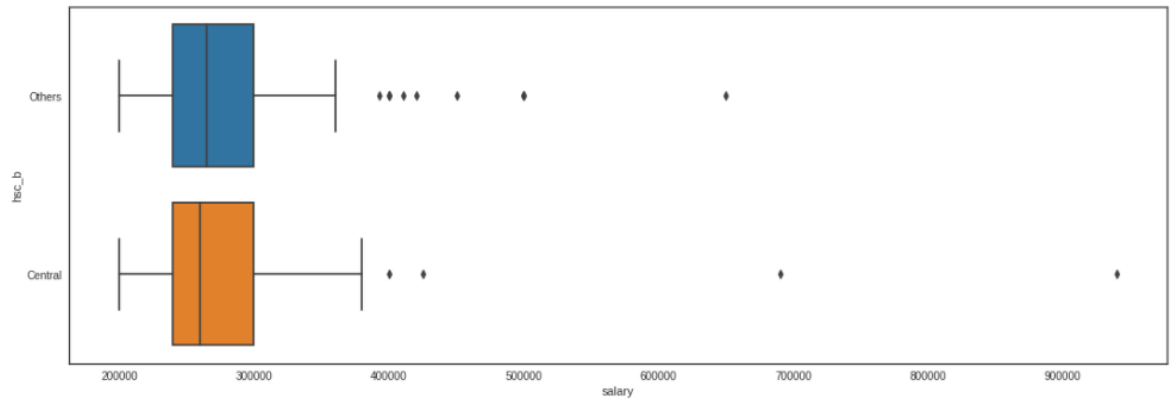
Education Board again, doesn't affect placement status much

```
In [107]: sns.countplot("hsc_s", hue="status", data=data)
plt.show()
```



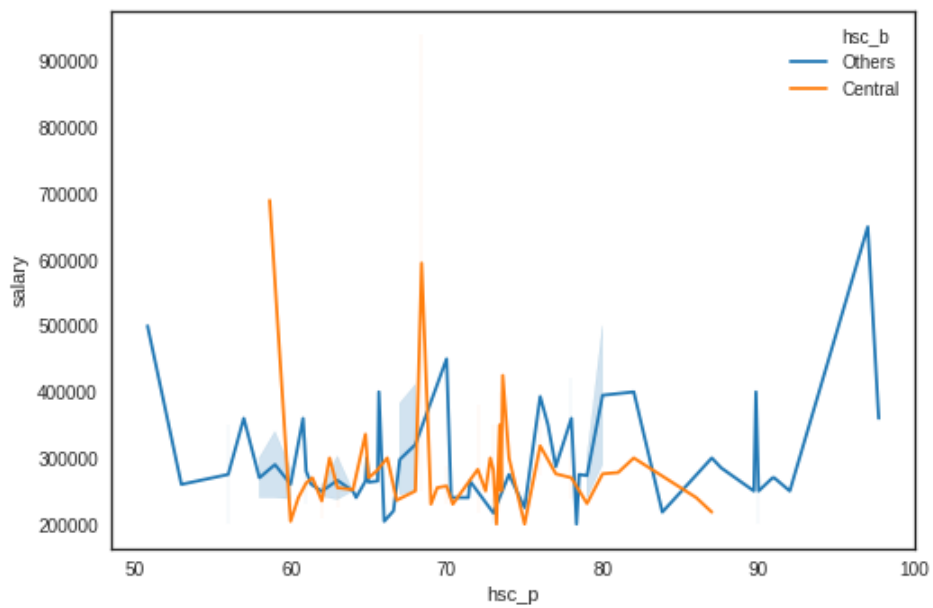
- We have very less students with Arts specialization.
- Around 2:1 placed:unplaced ratio for both Science and Commerce students

```
In [108]: plt.figure(figsize=(18,6))
sns.boxplot("salary", "hsc_b", data=data)
plt.show()
```



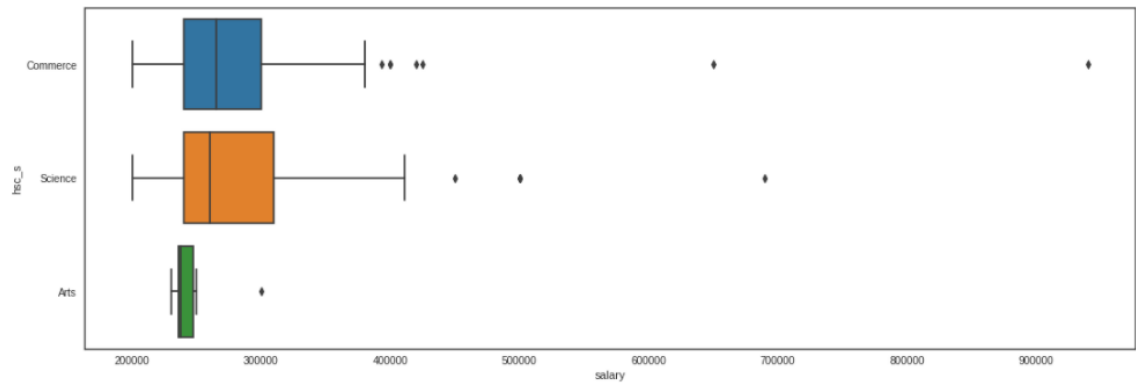
- Outliers on both, board doesn't affect getting highly paid jobs. Highest paid job was obtained by student from Central Board though.

```
In [109]: sns.lineplot("hsc_p", "salary", hue="hsc_b", data=data)
plt.show()
```



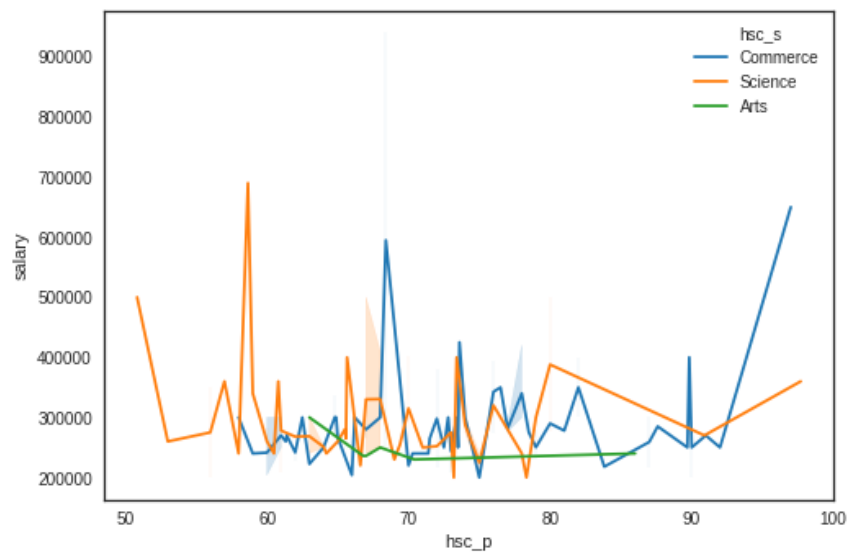
- High salary from both Central and Other.
- High salary for both high and low percentage.
- Thus, both these feature doesnot affect salary.

```
In [110]: plt.figure(figsize=(18,6))
sns.boxplot("salary", "hsc_s", data=data)
plt.show()
```



- We can't really say for sure due to only few samples of students with Arts Major, but they aren't getting good salaries.
- Commerce students have slightly better placement status.

```
In [111]: sns.lineplot("hsc_p", "salary", hue="hsc_s", data=data)
plt.show()
```

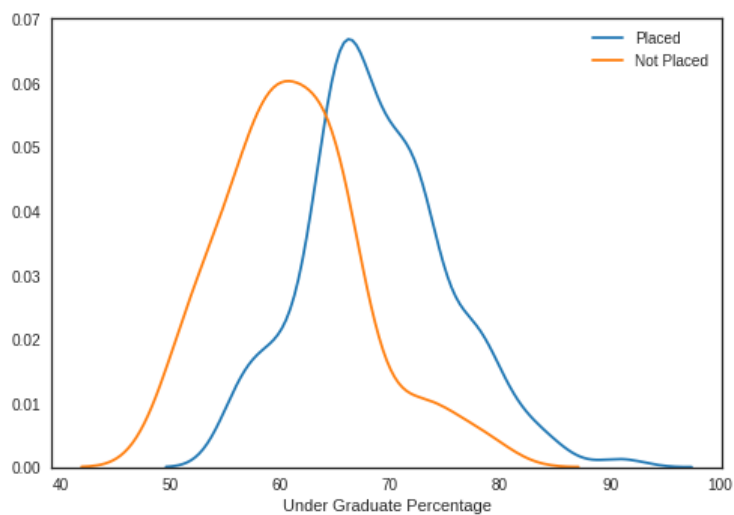


- **Student with Art Specialization surprisingly have comparatively low salary**

**Feature: degree\_p (Degree Percentage), degree\_t (Under Graduation Degree Field)**

**Does Under Graduate affect placements?**

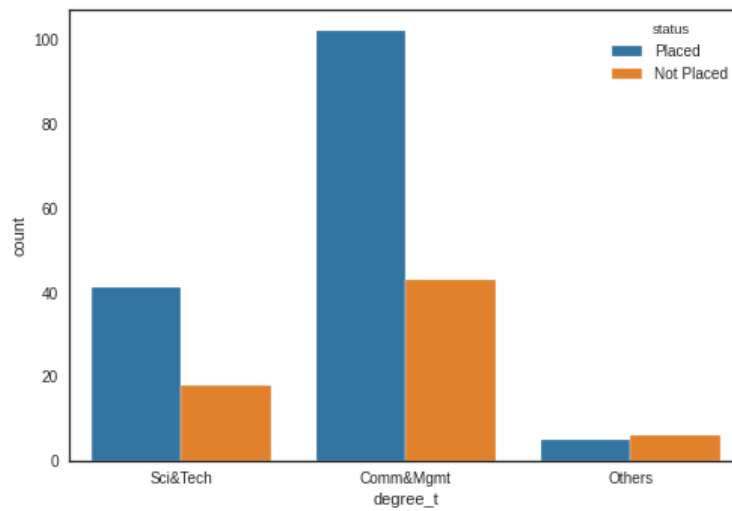
```
In [112]: #Kernel-Density Plot
sns.kdeplot(data.degree_p[ data.status=="Placed"])
sns.kdeplot(data.degree_p[ data.status=="Not Placed"])
plt.legend(["Placed", "Not Placed"])
plt.xlabel("Under Graduate Percentage")
plt.show()
```



- Overlap here too. But More placements for percentage above 65.
- UG Percentage least 50% to get placement

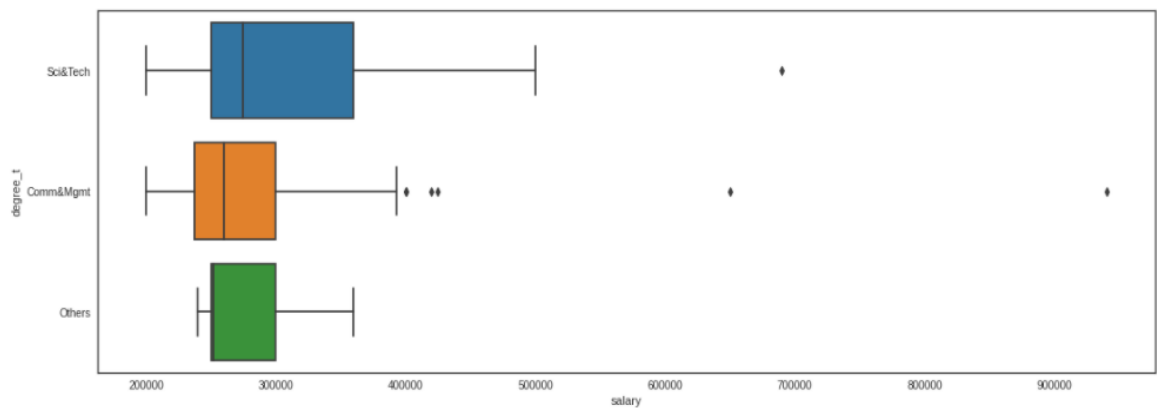


```
In [113]: sns.countplot("degree_t", hue="status", data=data)
plt.show()
```



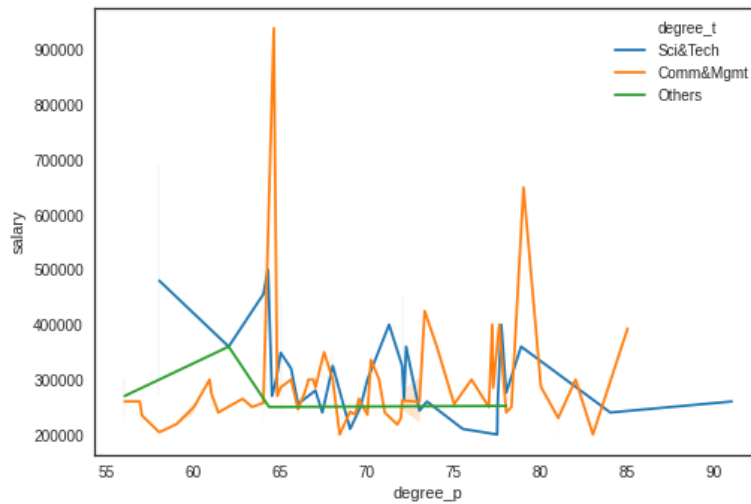
- We have very less students with "Other". We cant make decision from few cases.
- Around 2:1 placed:unplaced ratio for both Science and Commerce students

```
In [114]: plt.figure(figsize =(18,6))
sns.boxplot("salary", "degree_t", data=data)
plt.show()
```



- Science&Tech students getting more salary on average
  - Management stidents are getting more highly paid dream jobs.
-

```
In [115]: sns.lineplot("degree_p", "salary", hue="degree_t", data=data)
plt.show()
```

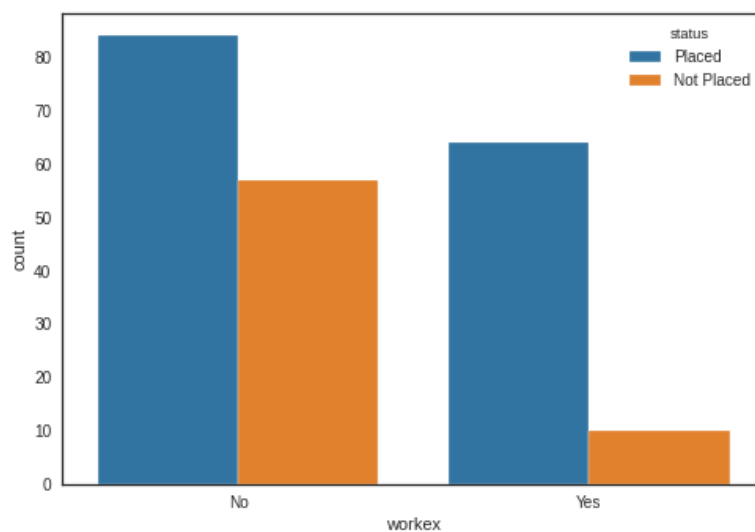


- Percentage does not seem to affect salary.
- Commerce&Mgmt students occasionally get dream placements with high salary

## Feature: Workex (Work Experience)

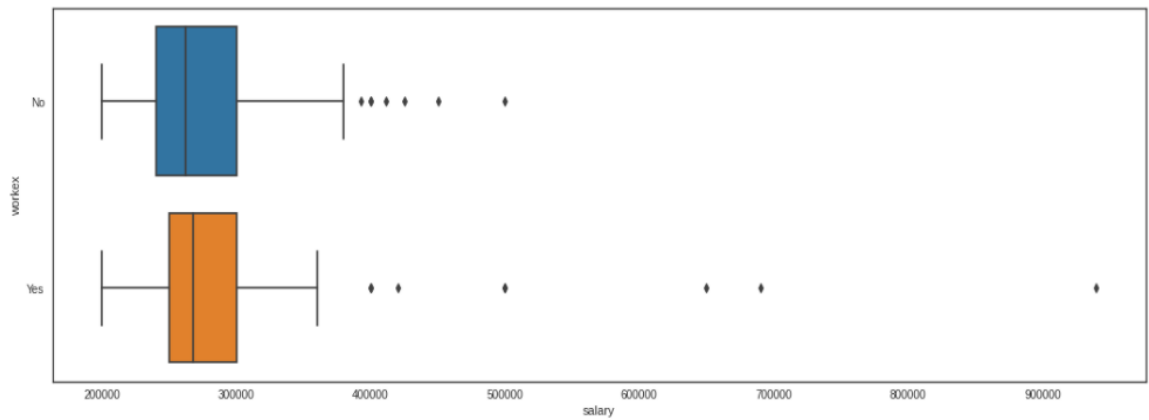
Does Work Experience affect placements?

```
In [116]: sns.countplot("workex", hue="status", data=data)
plt.show()
```



- **This affects Placement.** Very few students with work experience not getting placed!

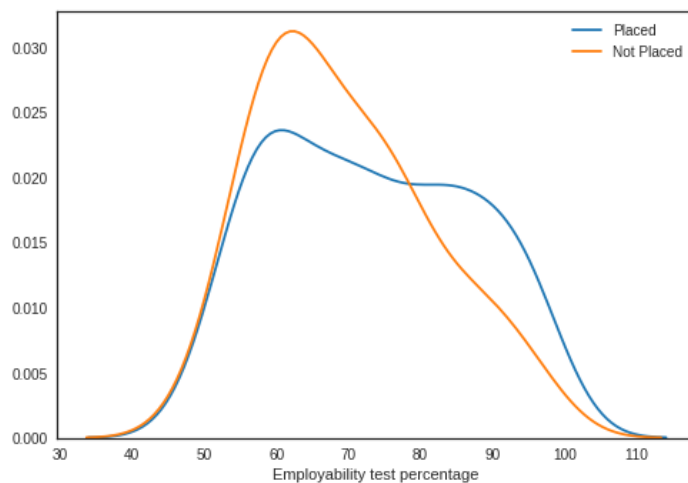
```
In [117]: plt.figure(figsize=(18,6))
sns.boxplot("salary", "workex", data=data)
plt.show()
```



- Outliers (High salary than average) on both end but **students with experience getting dream jobs**
- Average salary as well as base salary high for students with work experience

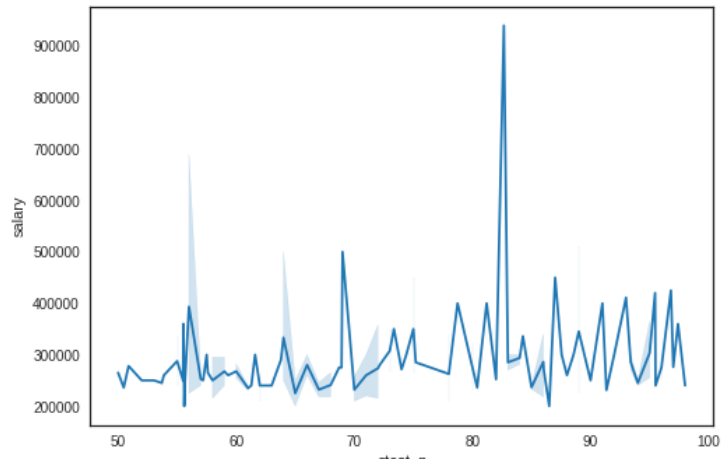
### Feature: etest\_p (Employability test percentage)

```
In [118]: #Kernel-Density Plot
sns.kdeplot(data.etest_p[ data.status=="Placed"])
sns.kdeplot(data.etest_p[ data.status=="Not Placed"])
plt.legend(["Placed", "Not Placed"])
plt.xlabel("Employability test percentage")
plt.show()
```



- High overlap -> It does not affect placement status much
- More "Not Placed" on percentage 50-70 range and more placed on 80% percentage range

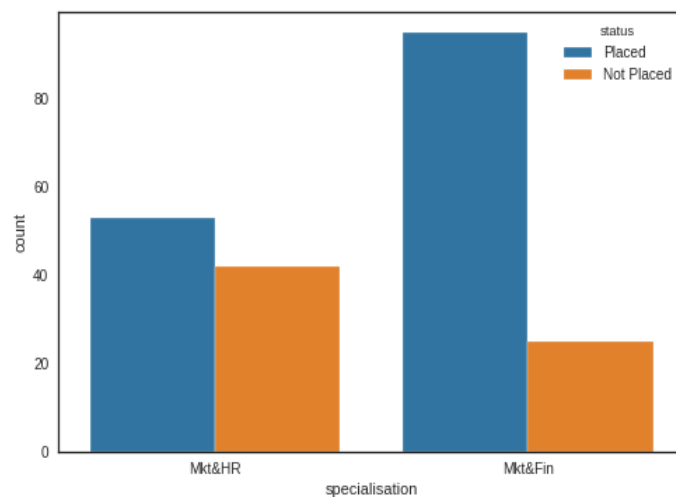
```
In [119]: sns.lineplot("etest_p", "salary", data=data)
plt.show()
```



This feature surprisingly does not affect placements and salary much

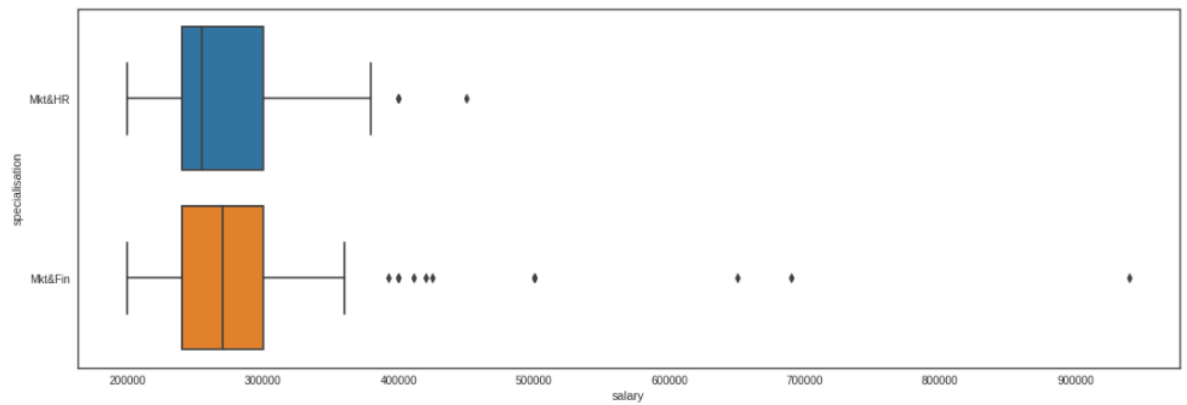
### Feature: Specialisation (Post Graduate Specialization)

```
In [120]: sns.countplot("specialisation", hue="status", data=data)
plt.show()
```



- This feature affects Placement status.
- Comparatively very low not-placed students in Mkt&Fin Section

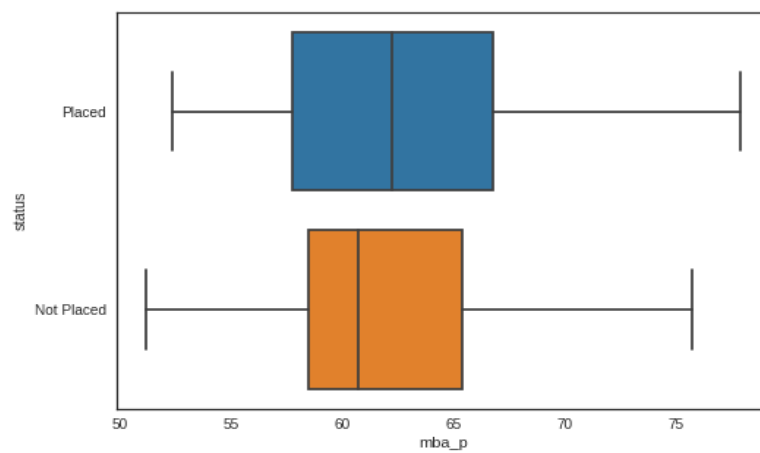
```
In [121]: plt.figure(figsize=(18,6))
sns.boxplot("salary", "specialisation", data=data)
plt.show()
```



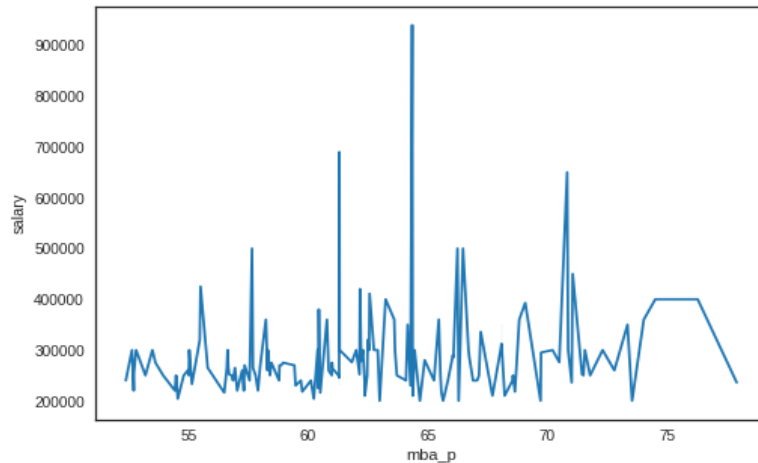
- *\*More Highly Paid Jobs for Mkt&Fin students \**

#### Does MBA Percentage affect placements?

```
In [122]: sns.boxplot("mba_p", "status", data=data)
plt.show()
```



```
In [123]: sns.lineplot("mba_p", "salary", data=data)
plt.show()
```



MBA Percentage also deos not affect salary much

## Feature Selection

Using Only following features (Ignoring Board of Education -> they didn't seem to have much effect)

- Gender
- Secondary Education percentage
- Higher Secondary Education Percentage
- Specialization in Higher Secondary Education
- Undergraduate Degree Percentage
- Under Graduation Degree Field
- Work Experience
- Employability test percentage
- Specialization
- MBA Percentage

Will compute feature importance later on.

## Data Pre-Processing

```
In [124]: data.drop(['ssc_b', 'hsc_b'], axis=1, inplace=True)
```

## Feature Encoding

```
In [125]: data.dtypes
# We have to encode gender, hsc_s, degree_t, workex, specialisation and status
```

```
Out[125]: gender          object
ssc_p          float64
hsc_p          float64
hsc_s          object
degree_p       float64
degree_t       object
workex         object
etest_p       float64
specialisation  object
mba_p          float64
status         object
salary         float64
dtype: object
```

```
In [126]: data["gender"] = data.gender.map({"M":0, "F":1})
data["hsc_s"] = data.hsc_s.map({"Commerce":0, "Science":1, "Arts":2})
data["degree_t"] = data.degree_t.map({"Comm&Mgmt":0, "Sci&Tech":1, "Others":2})
data["workex"] = data.workex.map({"No":0, "Yes":1})
data["status"] = data.status.map({"Not Placed":0, "Placed":1})
data["specialisation"] = data.specialisation.map({"Mkt&HR":0, "Mkt&Fin":1})
```

## Problem Statement

- Predicting If Students gets placed or not (Binary Classification Problem)
- Predicting Salary of Student (Regression Problem)

```
In [127]: # Lets make a copy of data, before we proceed with specific problems
data_clf = data.copy()
data_reg = data.copy()
```

## Binary Classification Problem

### Decision Tree Based Models

Using Decision Tree based Algorithm does not require feature scaling, and works great also in presence of categorical columns without ONE\_HOT Encoding

```
In [128]: # Library imports
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

## Dropping Salary Feature

Filling 0s for salary of students who didn't get placements would be bad idea as it would mean student gets placement if he earns salary.

```
In [129]: # Separating Features and Target
X = data_clf[['gender', 'ssc_p', 'hsc_p', 'hsc_s', 'degree_p', 'degree_t', 'workex', 'etest_p', 'specialisation', 'mba_p']]
y = data_clf['status']
```

```
In [130]: #Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [131]: dtree = DecisionTreeClassifier(criterion='entropy')
dtree.fit(X_train, y_train)
y_pred = dtree.predict(X_test)
```

```
In [132]: accuracy_score(y_test, y_pred)
```

```
Out[132]: 0.8307692307692308
```

```
In [133]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.65	0.76	0.70	17
1	0.91	0.85	0.88	48
accuracy			0.83	65
macro avg	0.78	0.81	0.79	65
weighted avg	0.84	0.83	0.83	65

```
In [134]: #Using Random Forest Algorithm
random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, y_train)
y_pred = random_forest.predict(X_test)
```

```
In [135]: accuracy_score(y_test, y_pred)
```

```
Out[135]: 0.9230769230769231
```

```
In [136]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.82	0.85	17
1	0.94	0.96	0.95	48
accuracy			0.92	65
macro avg	0.91	0.89	0.90	65

---



weighted avg      0.92      0.92      0.92      65

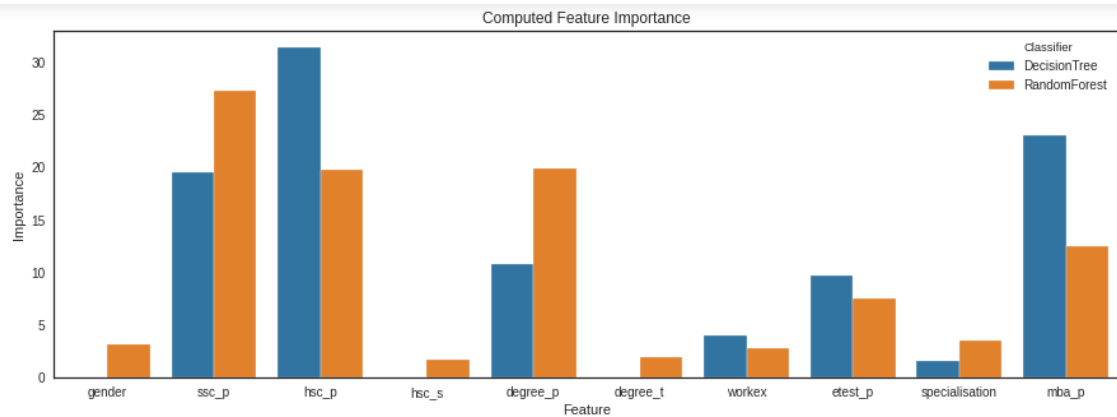
## Feature Importance (Percentage)

Tree based algorithms can be used to compute feature importance

Checking feature importance obtained from these:

```
In [137]: rows = list(X.columns)
imp = pd.DataFrame(np.zeros(6*len(rows)).reshape(2*len(rows), 3))
imp.columns = ["Classifier", "Feature", "Importance"]
#Add Rows
for index in range(0, 2*len(rows), 2):
    imp.iloc[index] = ["DecisionTree", rows[index//2], (100*dtree.feature_importances_[index//2])]
    imp.iloc[index + 1] = ["RandomForest", rows[index//2], (100*random_forest.feature_importances_[index//2])]

In [138]: plt.figure(figsize=(15,5))
sns.barplot("Feature", "Importance", hue="Classifier", data=imp)
plt.title("Computed Feature Importance")
plt.show()
```



hsc\_s -> Specialization in Higher Secondary Education

degree\_t -> Under Graduation(Degree type)- Field of degree education

specialisation -> Post Graduation(MBA)- Specialization

**Field of study does not seem to affect much**

Optionally we can remove these least important features and re-classify data.

## Binary Classification with Logistic Regression

### One Hot Encoding

Encoding Categorical Features

```
In [139]: # Separating Features and Target
X = data_clf[['gender', 'ssc_p', 'hsc_p', 'hsc_s', 'degree_p', 'degree_t', 'workex', 'etest_p', 'specialisation', 'mba_p']]
y = data_clf['status']
#Reverse Mapping and making Categorical
X["gender"] = pd.Categorical(X.gender.map({0:"M",1:"F"}))
X["hsc_s"] = pd.Categorical(X.hsc_s.map({0:"Commerce",1:"Science",2:"Arts"}))
X["degree_t"] = pd.Categorical(X.degree_t.map({0:"Comm&Mgmt",1:"Sci&Tech",2:"Others"}))
X["workex"] = pd.Categorical(X.workex.map({0:"No",1:"Yes"}))
X["specialisation"] = pd.Categorical(X.specialisation.map({0:"Mkt&HR",1:"Mkt&Fin"}))

In [140]: #One-Hot Encoding
X = pd.get_dummies(X)
column_names = X.columns.to_list()
```

### Feature Scaling

- Percentages are on scale 0-100
- Categorical Features are on range 0-1 (By one hot encoding)
- High Scale for Salary -> Salary is heavily skewed too -> SkLearn has RobustScaler which might work well here

Scaling Everything between 0 and 1 (This wont affect one-hot encoded values)

```
In [141]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

In [142]: #Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3)

In [143]: from sklearn.linear_model import LogisticRegression

In [144]: logistic_reg = LogisticRegression()
logistic_reg.fit(X_train, y_train)
y_pred = logistic_reg.predict(X_test)

In [145]: accuracy_score(y_test, y_pred)

Out[145]: 0.8
```

```
In [146]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.75	0.65	0.70	23
1	0.82	0.88	0.85	42
accuracy			0.80	65
macro avg	0.79	0.77	0.77	65
weighted avg	0.80	0.80	0.80	65

### Computing Feature importance by Mean Decrease Accuracy (MDA)

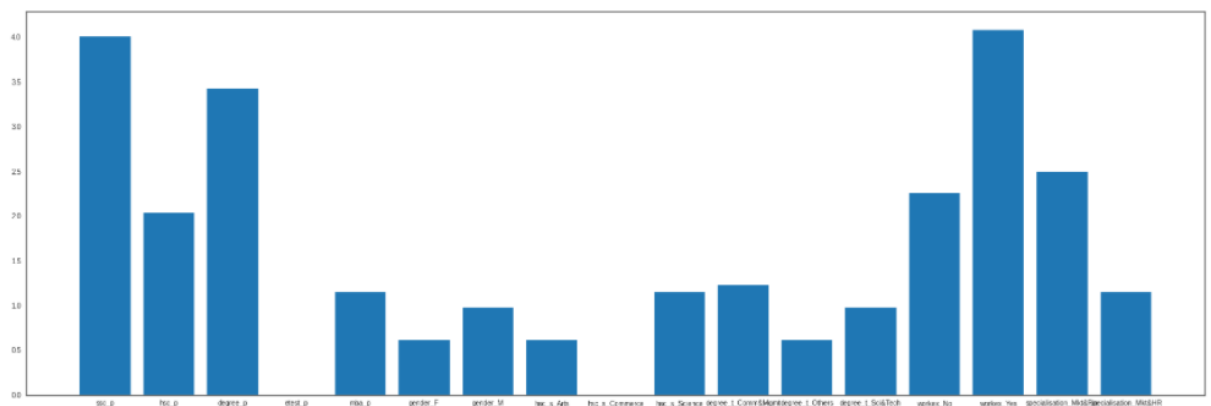
Since Logistic Regression performed well, Lets run another method for determining feature importance here.

```
In [147]: import eli5
from eli5.sklearn import PermutationImportance
perm = PermutationImportance(logistic_reg).fit(X_test, y_test)
eli5.show_weights(perm)
```

```
Out[147]:
```

Weight	Feature
0.1231 ± 0.0802	x0
0.0492 ± 0.0685	x2
0.0431 ± 0.0408	x1
0.0400 ± 0.0816	x14
0.0154 ± 0.0195	x12
0.0031 ± 0.0452	x13
0 ± 0.0000	x3
0 ± 0.0000	x8
-0.0031 ± 0.0123	x7
-0.0031 ± 0.0123	x11
-0.0092 ± 0.0246	x10
-0.0123 ± 0.0123	x5
-0.0123 ± 0.0230	x4
-0.0154 ± 0.0195	x6
-0.0185 ± 0.0230	x9
-0.0185 ± 0.0230	x16
-0.0215 ± 0.0500	x15

```
In [148]: plt.figure(figsize=(30, 10))
plt.bar(column_names , perm.feature_importances_std_ * 100)
plt.show()
```



**From Feature Importance of Tree-based Algorithms and MDA we can conclude that:**

- Academic performance affects placement (All percentages had importance)
- Work Experience Effects Placement
- Gender and Specialization in Commerce (in higher-secondary and undergraduate) also has effect on placements.

## Prediction of Salary (Regression Analysis)

```
In [149]: from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from sklearn.metrics import mean_absolute_error, r2_score
```

### Data Preprocessing

```
In [150]: #dropping NaNs (in Salary)
data_reg.dropna(inplace=True)
#dropping Status = "Placed" column
data_reg.drop("status", axis=1, inplace=True)
```

```
In [151]: data_reg.head()
```

```
Out[151]:
```

	gender	ssc_p	hsc_p	hsc_s	degree_p	degree_t	workex	etest_p	specialisation	mba_p	salary
0	0	67.00	91.00	0	58.00	1	0	55.0	0	58.80	270000.0
1	0	79.33	78.33	1	77.48	1	1	86.5	1	66.28	200000.0
2	0	65.00	68.00	2	64.00	0	0	75.0	1	57.80	250000.0
4	0	85.80	73.60	0	73.30	0	0	96.8	1	55.50	425000.0
7	0	82.00	64.00	1	66.00	1	1	67.0	1	62.14	252000.0

```
In [152]: #Seperating Dependent and Independent Variables
y = data_reg["salary"] #Dependent Variable
X = data_reg.drop("salary", axis=1)
column_names = X.columns.values
```

```
In [153]: #Scalizing between 0-1 (Normalization)
X_scaled = MinMaxScaler().fit_transform(X)
```

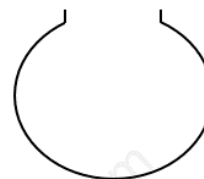
### Feature Selection

\*\* Not all features are significant. Thus, let's perform a feature selection procedure\*\*

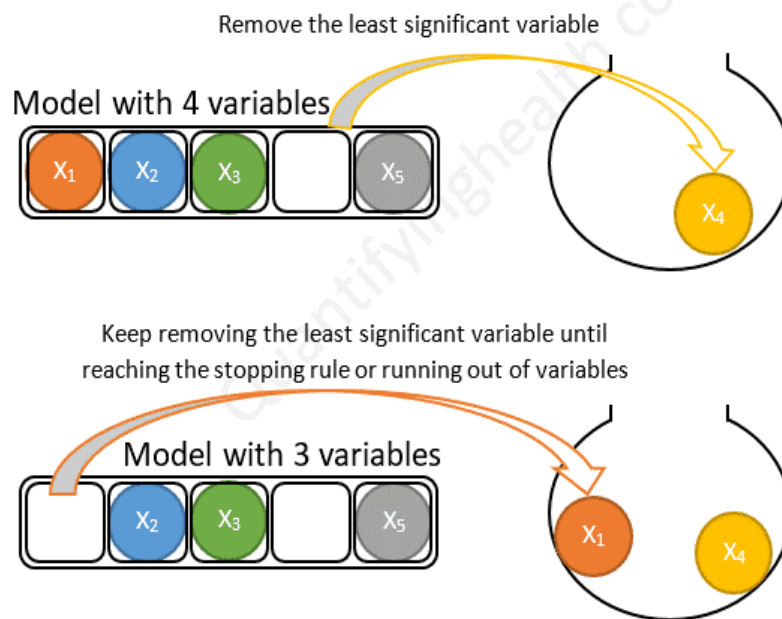
#### Backward stepwise selection example with 5 variables:

Start with a model that contains all the variables

Full Model



Remove the least significant variable



### Determining Least Significant Variable

The least significant variable is a variable which:

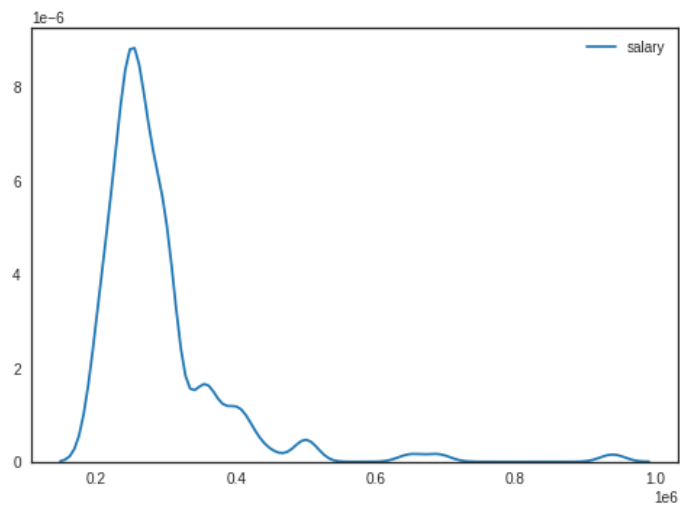
- has the highest p-value
- Removing it reduces  $R^2$  to lowest value compared to other features
- Removing it has least increment in residuals-sum-of-squares (RSS)

### Outliers' Removal

Feature Selection cannot perform well in presence of outliers. Let's identify and remove outliers before proceeding

```
In [154]: #PDF of Salary
sns.kdeplot(y)
plt.show()
```

```
In [154]: #PDF of Salary
sns.kdeplot(y)
plt.show()
```



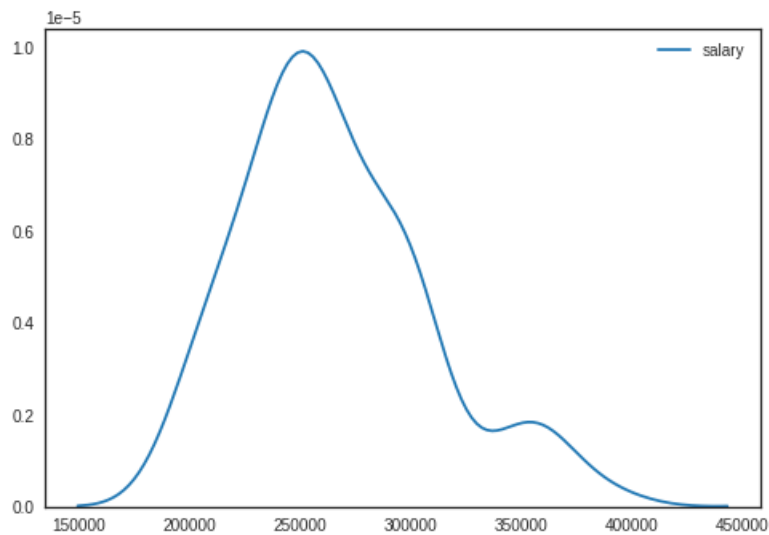
It is clear that very few students have salary greater than 400,000 (hence outliers)

```
In [155]: #Selecting outliers
y[y > 400000]
# 9 records
```

```
Out[155]: 4      425000.0
          39      411000.0
          53      450000.0
          77      500000.0
          95      420000.0
          119     940000.0
          150     690000.0
          163     500000.0
          174     500000.0
          177     650000.0
          Name: salary, dtype: float64
```

```
In [156]: #Removing these Records from data
X_scaled = X_scaled[y < 400000]
y = y[y < 400000]
```

```
In [157]: #PDF of Salary without outliers. Still skewed though
sns.kdeplot(y)
plt.show()
```

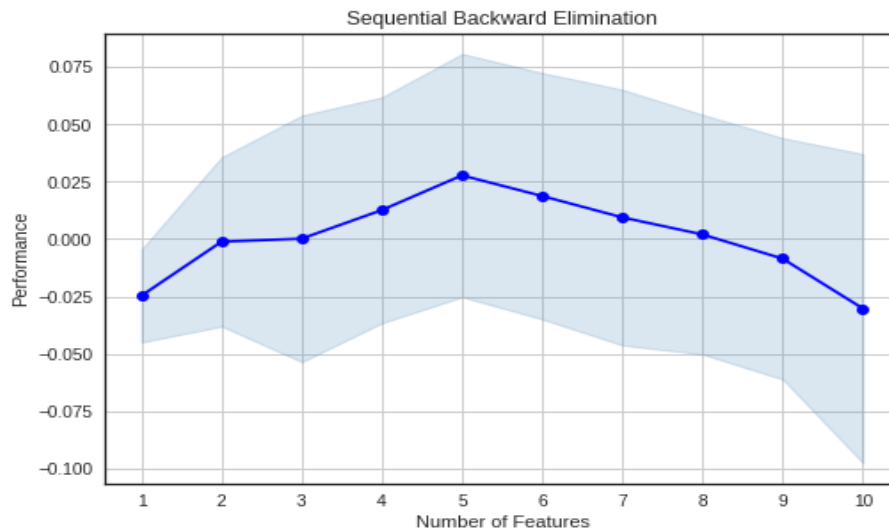


## 1. Determining Least Significant Variable by R2 Score

```
In [158]: from mlxtend.feature_selection import SequentialFeatureSelector as SFS
          from mlxtend.plotting import plot SequentialFeatureSelection as plot_sfs
```

```
In [159]: linreg = LinearRegression()
          sfs = SFS(linreg, k_features=1, forward=False, scoring='r2', cv=10)
          sfs = sfs.fit(X_scaled, y)
          fig = plot_sfs(sfs.get_metric_dict(), kind='std_err')

          plt.title('Sequential Backward Elimination')
          plt.grid()
          plt.show()
          #From Plot its clear that, many features actually decrease the performance
```



```
In [160]: # Lets see the top 5 most significant features
top_n = 5
sfs.get_metric_dict()[top_n]
```

```
Out[160]: {'feature_idx': (0, 3, 5, 7, 9),
'cv_scores': array([-0.09997564, -0.11551795,  0.14652782,  0.19241391, -0.19535134,
-0.13235138, -0.03896556,  0.3116134 ,  0.13836643,  0.07020936]),
'avg_score': 0.027696903219017056,
'feature_names': ('0', '3', '5', '7', '9'),
'ci_bound': 0.11802751969012426,
'std_dev': 0.15891404557580263,
'std_err': 0.05297134852526754}
```

```
In [161]: #Top N Features
top_n_indices = list(sfs.get_metric_dict()[top_n]['feature_idx'])
print(f"Most Significant {top_n} Features:")
for col in column_names[top_n_indices]:
    print(col)
```

```
Most Significant 5 Features:
gender
hsc_s
degree_t
etest_p
mba_p
```

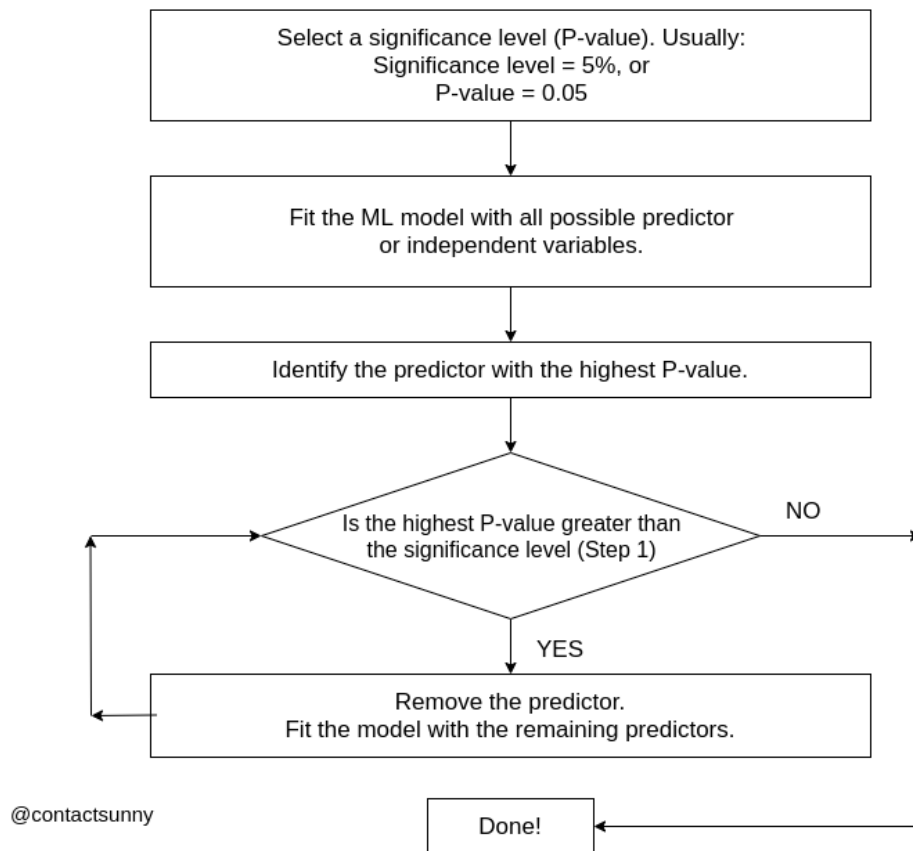
```
In [162]: #Select these Features only
X_selected = X_scaled[:,top_n_indices]
lin_reg = LinearRegression()
lin_reg.fit(X_selected, y)
y_pred = lin_reg.predict(X_selected)
print(f"R2 Score: {r2_score(y, y_pred)}")
print(f"MAE: {mean_absolute_error(y, y_pred)}")
```

```
R2 Score: 0.1101660718969637
MAE: 30630.128295211573
```

This is the best I could do with Linear Regression



## Determining Least Significant Variable by P-Value



```
In [163]: #Converting to DF for as column names gives readability
X_scaled = pd.DataFrame(X_scaled, columns=column_names)
y = y.values

# We must add a constants 1s for intercept before doing Linear Regression with statsmodel
X_scaled = sm.add_constant(X_scaled)
X_scaled.head()
#Constants 1 added for intercept term
```

```
Out[163]:
```

	const	gender	ssc_p	hsc_p	hsc_s	degree_p	degree_t	workex	etest_p	specialisation	mba_p
0	1.0	0.0	0.445545	0.857051	0.0	0.057143	0.5	0.0	0.104167	0.0	0.251666
1	1.0	0.0	0.750743	0.586729	0.5	0.613714	0.5	1.0	0.760417	1.0	0.544884
2	1.0	0.0	0.396040	0.366332	1.0	0.228571	0.0	0.0	0.520833	1.0	0.212466
3	1.0	0.0	0.816832	0.280990	0.5	0.285714	0.5	1.0	0.354167	1.0	0.382595
4	1.0	0.0	0.594059	0.601024	0.0	0.457143	0.0	0.0	0.861250	1.0	0.349275

# OLS Regression Results

Dep. Variable:	y	R-squared:	0.123
Model:	OLS	Adj. R-squared:	0.052
Method:	Least Squares	F-statistic:	1.722
Date:	Sat, 05 Jun 2021	Prob (F-statistic):	0.0829
Time:	05:40:18	Log-Likelihood:	-1608.4
No. Observations:	134	AIC:	3239.
Df Residuals:	123	BIC:	3271.
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.625e+05	1.28e+04	20.498	0.000	2.37e+05	2.88e+05
gender	-1.784e+04	8299.998	-2.149	0.034	-3.43e+04	-1406.775
ssc_p	-116.6148	2.04e+04	-0.006	0.995	-4.04e+04	4.02e+04
hsc_p	-1.842e+04	2.13e+04	-0.864	0.389	-6.06e+04	2.38e+04
hsc_s	-2.775e+04	1.58e+04	-1.761	0.081	-5.9e+04	3444.983
degree_p	-9885.6991	2.25e+04	-0.438	0.662	-5.45e+04	3.47e+04
degree_t	3.947e+04	1.69e+04	2.340	0.021	6077.584	7.29e+04
workex	-7748.2212	7673.070	-1.010	0.315	-2.29e+04	7440.151
etest_p	1.839e+04	1.43e+04	1.286	0.201	-9906.447	4.67e+04
specialisation	2457.2424	8013.710	0.307	0.760	-1.34e+04	1.83e+04
mba_p	3.704e+04	2.11e+04	1.756	0.082	-4717.648	7.88e+04

Omnibus:	10.852	Durbin-Watson:	1.965
Prob(Omnibus):	0.004	Jarque-Bera (JB):	11.041
Skew:	0.661	Prob(JB):	0.00400
Kurtosis:	3.477	Cond. No.	12.9

## RESULT

### OLS Regression Results

Dep. Variable:	y	R-squared:	0.123
Model:	OLS	Adj. R-squared:	0.052
Method:	Least Squares	F-statistic:	1.722
Date:	Sat, 05 Jun 2021	Prob (F-statistic):	0.0829
Time:	05:40:18	Log-Likelihood:	-1608.4
No. Observations:	134	AIC:	3239.
Df Residuals:	123	BIC:	3271.
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.625e+05	1.28e+04	20.498	0.000	2.37e+05	2.88e+05
gender	-1.784e+04	8299.998	-2.149	0.034	-3.43e+04	-1406.775
ssc_p	-116.6148	2.04e+04	-0.006	0.995	-4.04e+04	4.02e+04
hsc_p	-1.842e+04	2.13e+04	-0.864	0.389	-6.06e+04	2.38e+04
hsc_s	-2.775e+04	1.58e+04	-1.761	0.081	-5.9e+04	3444.983
degree_p	-9885.6991	2.25e+04	-0.438	0.662	-5.45e+04	3.47e+04
degree_t	3.947e+04	1.69e+04	2.340	0.021	6077.584	7.29e+04
workex	-7748.2212	7673.070	-1.010	0.315	-2.29e+04	7440.151
etest_p	1.839e+04	1.43e+04	1.286	0.201	-9906.447	4.67e+04
specialisation	2457.2424	8013.710	0.307	0.760	-1.34e+04	1.83e+04
mba_p	3.704e+04	2.11e+04	1.756	0.082	-4717.648	7.88e+04

Omnibus:	10.852	Durbin-Watson:	1.965
Prob(Omnibus):	0.004	Jarque-Bera (JB):	11.041
Skew:	0.661	Prob(JB):	0.00400
Kurtosis:	3.477	Cond. No.	12.9

```
In [165]: # Identify max P-value (P>|t|) column
# Feature ssc_p has 0.995
#drop ssc_p
X_scaled = X_scaled.drop('ssc_p', axis=1)
model = sm.OLS(y, X_scaled)
results = model.fit()
results.summary()
```

Out[165]: OLS Regression Results

Dep. Variable:	y	R-squared:	0.123
Model:	OLS	Adj. R-squared:	0.059
Method:	Least Squares	F-statistic:	1.929
Date:	Sat, 05 Jun 2021	Prob (F-statistic):	0.0536
Time:	05:40:18	Log-Likelihood:	-1608.4
No. Observations:	134	AIC:	3237.
Df Residuals:	124	BIC:	3266.
Df Model:	9		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.625e+05	1.2e+04	21.888	0.000	2.39e+05	2.86e+05
gender	-1.784e+04	8177.285	-2.182	0.031	-3.4e+04	-1657.933
hsc_p	-1.845e+04	2.08e+04	-0.888	0.376	-5.96e+04	2.27e+04

Dep. Variable:	y	R-squared:	0.079			
Model:	OLS	Adj. R-squared:	0.057			
Method:	Least Squares	F-statistic:	3.703			
Date:	Sat, 05 Jun 2021	Prob (F-statistic):	0.0135			
Time:	05:40:18	Log-Likelihood:	-1611.7			
No. Observations:	134	AIC:	3231.			
Df Residuals:	130	BIC:	3243.			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	2.573e+05	7260.622	35.441	0.000	2.43e+05	2.72e+05
gender	-2.036e+04	7777.446	-2.618	0.010	-3.57e+04	-4973.494
degree_t	2.202e+04	1.33e+04	1.654	0.101	-4325.122	4.84e+04
mba_p	3.138e+04	1.67e+04	1.884	0.062	-1567.309	6.43e+04
Omnibus:	13.932	Durbin-Watson:	2.008			
Prob(Omnibus):	0.001	Jarque-Bera (JB):	14.967			
Skew:	0.773	Prob(JB):	0.000562			
Kurtosis:	3.539	Cond. No.	5.77			

## CONCLUSION

From the machine learning models we identify what are all the features that affect the salary.

Therefore the top 5 features that affect the salary are:

- gender -> Gender
- degree\_t -> Under Graduation(Degree type)- Field of degree education
- mba\_p -> MBA percentage
- hsc\_s -> Specialization in Higher Secondary Education
- etest\_p -> Employability test percentage

## REFERENCES

- <https://www.kaggle.com/albertonavaa/beginner-approach-to-campus-placement-prediction>
- [kaggle.com/cs155925/campus-recruitment-exploration](https://www.kaggle.com/cs155925/campus-recruitment-exploration)
- <https://www.kaggle.com/srikardornala/placement-dataset-regression-classification>
- <https://www.kaggle.com/nareshbhat/outlier-the-silent-killer>
- <https://www.kaggle.com/benroshan/you-re-hired-analysis-on-campus-recruitment-data>